

FPR Annotation Task — Briefing Sheet

What you are doing

You will read 225 social media posts and judge whether a specific word or phrase in each post is being used as hateful coded language or not. This task should take approximately 4–7 hours total. You do not need to complete it in one sitting.

The spreadsheet columns

The sheet has five visible columns:

row_id — A number from 1 to 225. This is just for reference so you can track your progress and ask questions about specific rows.

matched_surface_form — The specific word or phrase you are judging. This is a term that has been documented as coded hate speech (a "dogwhistle") in academic research. Your job is to judge whether it is being used that way in this particular post.

post_text — The full text of the social media post. Read the whole post, but focus your judgment on the **matched_surface_form** term specifically.

judgment — This is where you enter your answer. Type one of three options (see below). This is the only column you need to fill in.

annotator_notes — Optional. Use this to flag anything unusual about a post, explain an Unclear judgment, or note if you think the matched term does not appear in the post text. Leave blank if you have nothing to add.

The three judgment options

Enter exactly one of these three values in the **judgment** column:

Hateful — The term is being used with its hateful coded meaning in this post. The author appears to be expressing or amplifying bias against a group.

Not Hateful — The term is not being used hatefully in this post. This includes posts that quote or discuss the term, report on it neutrally, or use it in an entirely unrelated context.

Unclear — You genuinely cannot tell from the post alone whether the term is being used hatefully or not. Use this freely — it is a valid answer, not a cop-out.

What to focus on

Your judgment should be about **how the matched term is being used in this specific post**, not about whether the post is offensive overall. A post can contain slurs and still be "Not Hateful" for the specific

matched term if that term is being used neutrally. A post can be polite in tone and still be "Hateful" if the matched term is clearly being used as coded hate speech.

Ask yourself: **is the author using this word or phrase to express or amplify hatred toward a group, or not?**

Important notes

- Some posts contain slurs and offensive language. This is expected and unavoidable given the source material.
- Some posts are genuinely ambiguous. Use **Unclear** when uncertain — do not guess between Hateful and Not Hateful.
- If a post contains `<number>` or `<percent>`, these are placeholders for numbers that were anonymized in the original dataset. Ignore them.
- If you cannot find the matched surface form anywhere in the post text, note this in `annotator_notes` and mark the judgment as **Unclear**.
- Do not look up the post, the term, or any related context online. Judge only from what is in the post text.
- Aim for roughly 1–2 minutes per post. Some will be obvious in seconds; some will take longer.